

# TUTORBENCH: A BENCHMARK TO ASSESS TUTORING CAPABILITIES OF LARGE LANGUAGE MODELS

**Rakshith S Srinivasa, Zora Che, Chen Bo Calvin Zhang, Diego Mares, Ernesto Hernandez, Jayeon Park, Dean Lee, Guillermo Mangialardi, Charmaine Ng, Ed-Yeremai Hernandez Cardona, Anisha Gunjal, Yunzhong He, Bing Liu, Chen Xing**

Scale AI

## ABSTRACT

As students increasingly adopt large language models (LLMs) as learning aids, it is essential to evaluate their tutoring capabilities rigorously. In this paper, we introduce TUTORBENCH, a dataset and evaluation benchmark designed to evaluate the core tutoring skills of LLMs. The dataset comprises 1,490 prompts curated by human experts, focused on high-school and AP-level curricula. TUTORBENCH consists of examples drawn from three common tutoring tasks: (i) generating adaptive explanations tailored to a student’s confusion, (ii) providing actionable feedback on a student’s work, and (iii) promoting active learning through effective hint generation. TUTORBENCH is accompanied by a reliable and fine-grained automatic evaluation method that uses an LLM-judge and sample-specific rubrics. We evaluate 16 frontier LLMs on TUTORBENCH and present a detailed analysis of their performance and behavior. Our results show that none of the frontier LLMs achieve a score of greater than 56%, showing a large room for improvement. We find that LLMs still fall short in exhibiting the full range of tutoring skills needed to guide, diagnose, and support students’ reasoning effectively, with all the frontier models achieving less than 60% pass rate on rubric criteria related to these skills. We also find that different model families exhibit varied strengths and limitations: the Claude models outperform others in supporting active learning, while they trail behind them in the other two use cases. By releasing TUTORBENCH and an accompanying leaderboard, we provide a comprehensive and unsaturated benchmark to guide the development of the next-generation of AI tutors.

## 1 INTRODUCTION

Large language models (LLMs) are rapidly transforming the way students learn and access educational support (Handa et al., 2025; Ammari et al., 2025; Geraghty & Goldstein, 2024; Scarlatos et al., 2025). Tools like ChatGPT, Gemini, and Claude have already begun serving as on-demand tutors for millions of learners worldwide. To encourage adoption, many frontier model providers offer variants specializing in tutoring (OpenAI, 2025b) and extend free access to students (Google, 2025). It is clear that AI models promise scalable, personalized tutoring at low cost. However, their true impact is not yet fully understood: while many studies show improvements in learning outcomes (Vanzo et al., 2025), there are also reports that students perform worse when the tools are taken away (Bastani et al., 2025).

Given their rapid adoption, there is an urgent need for realistic and thorough ways to evaluate LLMs on their tutoring capabilities. Many existing LLM benchmarks focus on frontier capabilities such as advanced domain knowledge, reasoning, and problem-solving (Phan et al., 2025; Glazer et al., 2024; He et al., 2024). While these dimensions are crucial for understanding the limits of model performance, they overlook more nuanced human-centered capabilities necessary for tutoring, such as clear and thorough explanation, adaptivity to a learner’s needs, and support for active learning. Some recent efforts have begun to explore how to benchmark LLMs in tutoring contexts. However, they are limited in scope and easily saturated by current state-of-the-art (SOTA) LLMs. For example, authors in Maurya et al. (2025); Gupta et al. (2025) focus on text-only examples on one or two subjects like mathematics or computer science. This differs significantly from real-world usage, where students frequently upload pictures of their work and seek help on several subjects. SOTA

 **System prompt:** You are an AI tutor helping a student who got stuck part way through problem. Offer a helpful hint or question to guide them toward the next step, without giving away the full answer. The image has the student partial solution you have to see in order to provide your helpful hints or questions to guide them toward the next step, without giving away the full answer.

 **Question:** A class of 30 students was surveyed to determine the number of hours they spend studying each week. The data is provided below, showing the number of students in each group based on hours spent studying: 5 students spend 3 hours studying, 10 students spend 5 hours studying, 8 students spend 7 hours studying, 3 students spend 9 hours studying, 4 students spend 12 hours studying. Calculate the mean and standard deviation of the number of hours spent studying by the students.

 Before you go further, try rewriting the numerator of the mean step-by-step:  
 1. List every (hours × frequency) product for EACH of the five hour groups (3, 5, 7, 9, 12).  
 2. Add those five products together.  
 3. Divide that total by the full class size (30 students).

Does your current expression include the 9-hour and 12-hour groups? Once you have the correct mean, use it in  $\sigma^2 = \sum (x - \mu)^2 / N$  (remember to square the difference first, then multiply by the frequency, and finally sum over all groups).

$$\begin{aligned} \text{Mean} &= \frac{(3+5) + (5+10) + (7+8) + (9+3) + (12+4)}{30} \\ &= \frac{15 + 50 + 56 + 27 + 48}{30} \\ &= \frac{196}{30} \approx 6.53 \end{aligned}$$



standard deviation :  
 I know that it relate to the mean I calculated  
 but I don't remember the formula I thought  
 it could be  $\frac{\sqrt{\sum (x - \mu)^2}}{N}$ , but I'm not sure.

Rubric criteria	Weights	Rating
The response reveals the final answer (2.74) directly	-5	
The response should use clear formatting or structure (e.g., headers) to organize explanation into sections.	1	
⋮		
The response must explain that taking the square root of the variance gives the standard deviation in original units.	5	

Figure 1: An example from the TUTORBENCH dataset. Each sample includes a **system prompt** defining the tutoring goal (top left), a **question** along with the student’s partial work (top right), an **AI model’s response** (bottom left), and a set of **rubric criteria** for evaluation (bottom right). In this instance, the AI is prompted to provide a hint for a statistics problem.

LLMs such as GPT-4 achieve nearly perfect scores on both of these benchmarks according to their own reported results, indicating the need for more nuanced benchmarks that expose shortcomings in model behavior. Moreover, they lack reliable auto-judge methods to evaluate LLMs’ tutoring capability on open-ended and subjective tutoring tasks mentioned above.

To address the above gaps, we present **TUTORBENCH**<sup>1</sup>, a benchmark to evaluate the tutoring capabilities of LLMs. TUTORBENCH consists of 1490 conversations between a student-persona and a tutor-persona. The conversations cover 6 STEM subjects (Biology, Physics, Chemistry, Statistics, Calculus, and Computer Science) and focus on high school and AP curricula. The conversations are designed to reflect real-world usage by students and focus on three tutoring use cases: adaptive explanation generation, feedback and assessment, and active learning support. These use cases were chosen to test an LLM’s ability to calibrate its responses to individual students and support their learning journey rather than simply generating universal and standard solutions to questions. TUTORBENCH is also *multimodal*, with both text-only and image-based conversations. 828 out of the 1490 samples in this dataset contain images of handwritten or printed work by students. This reflects real-world usage by students, where uploading images of their work to models is common. We provide more details about the use cases in Section 2.1.

We design TUTORBENCH to be a *rubrics-based* evaluation to account for the open-ended nature of tutoring. Each sample in the dataset is accompanied by a set of sample-specific evaluation rubrics written by human experts that are self-contained, mutually exclusive, and verifiable. Together, these rubrics capture the requirements of a desirable response to the corresponding test sample. Model responses are graded by an LLM judge with a pass/fail rating on each rubric criterion.

We conduct data collection of TUTORBENCH with the help of human experts. The questions and the rubrics to each question were written by experts of the corresponding subject who hold a Bachelor’s or higher degree and have either tutoring or professional experience in the corresponding subject. In order to guarantee the difficulty level of TUTORBENCH, we then prompt 5 state-of-the-art LLMs (Gemini 2.5 Pro Google DeepMind (2025), Claude 3.7 Sonnet Anthropic (2025a), o3 OpenAI (2025a), DeepSeek-R1 DeepSeek (2025) and Llama 4 Maverick Meta AI (2025)) to respond to

<sup>1</sup>A sample subset of TUTORBENCH can be found at [https://huggingface.co/datasets/ScaleAI/tutorbench\\_sample](https://huggingface.co/datasets/ScaleAI/tutorbench_sample). The full dataset will be released soon.

each collected test sample. We retain only conversations that result in a score of less than 50% for at least three of the five models. This results in a challenging benchmark, with the best performing model attaining a score of 55.65%.

Using the samples in TUTORBENCH, we evaluate and analyze 16 frontier LLMs. We report the overall scores of the models in Section 4.1. Among all the models tested, Gemini 2.5 Pro and GPT-5 achieve the best performance, with 55.65% and 55.33% final scores respectively (100% being the maximum final score). The Claude model series falls behind the Gemini 2.5 Pro and GPT model series with non-negligible performance gaps, with Claude Opus 4.1 (Thinking) achieving a final score of 50.78%. We observe that models achieve a score of only 47.16% on the use case of *adaptive explanation generation*, indicating that frontier LLMs still struggle to generate effective personalized responses. Further, they also achieve less than 55% on the other two use cases. Interestingly, models of the Claude series outperform others in the “*active learning support*” use case, while trailing in the overall performance. This shows their unique focus on model training for tutoring purposes. More detailed analysis and observations with respect to the use cases and evaluation dimensions can be found in Section 4.2 and Section 4.3.

Overall, TUTORBENCH is a comprehensive, multimodal, and challenging dataset to evaluate frontier LLMs for their tutoring capabilities. Our contributions are as follows:

- A dataset with 1490 realistic conversations between students and tutors that cover 3 use cases on 6 STEM subjects. 828 of the conversations are multimodal, reflecting real-world usage by students.
- The conversations are accompanied by specific rubric criteria that allow for reliable automatic model performance evaluation.
- We provide an evaluation of 16 frontier models on TUTORBENCH. We also analyze models on finer aspects of tutoring and identify areas of improvement.
- To further facilitate tracking model performance and research, we release a leaderboard and a sample subset of TUTORBENCH, and we will release the full dataset soon.

## 2 BENCHMARK DESIGN

We design TUTORBENCH to cover a broad spectrum of common, real-world tutoring use cases, including both text-only and multimodal questions on 6 STEM subjects. Use cases in TUTORBENCH focus on assessing human-centered capabilities necessary for tutoring, such as explanation, guidance, and adaptivity to a learner’s needs. We also design a reliable rubric-based eval to automatically assess LLM capabilities on such subjective matters.

### 2.1 TUTORING USE CASES

As large language models (LLMs) become increasingly integrated into educational workflows, it is important to rigorously evaluate their effectiveness in different tutoring scenarios. TUTORBENCH captures three core use cases that reflect essential tutoring behaviors: adaptive explanation, assessment and feedback, and active learning support. We elaborate on these 3 use cases below.

**Use Case 1: Adaptive Explanation Generation.** Personalized instructions can help unlock a student’s potential. One of the most impactful qualities of a good human tutor is their ability to adapt explanations based on a student’s current understanding and knowledge gaps. The key skill required here is to identify core misconceptions and errors in a student’s understanding and address them in a way that is easy to understand based on the current status of the student.

To evaluate this capability, we design a multi-turn interaction setup consisting of: an initial question written in a *student-persona*, an initial explanation written in a *tutor-persona*, and a follow-up question written in a student-persona. During evaluation, an LLM is presented with this triplet, along with a system prompt/instruction to generate an explanation that directly addresses the knowledge gap or misunderstanding exhibited in the follow-up question by the student. An example conversation is shown in Fig. 2. Model responses are evaluated on whether they can recognize the specific context and background implied by the student’s follow-up, and whether they can produce a helpful, focused, and to-the-point response tailored to the student’s current understanding.

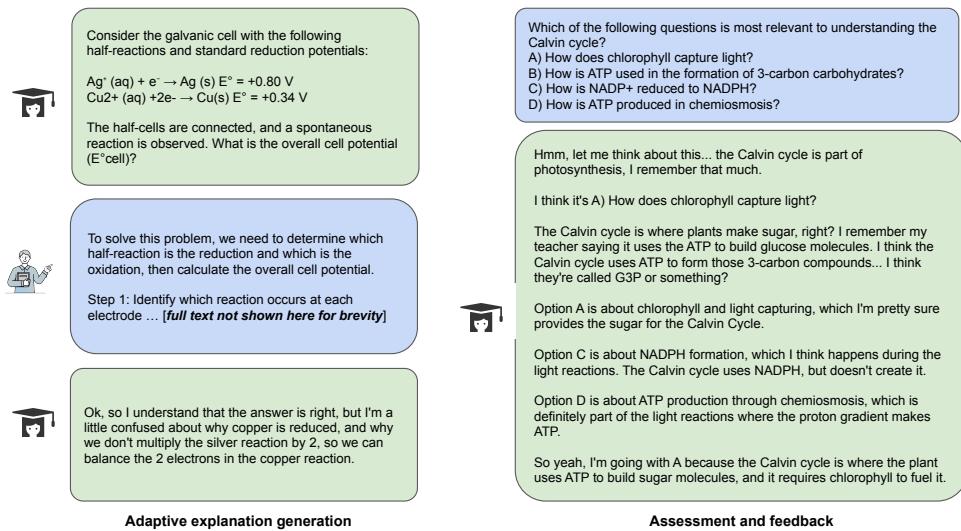


Figure 2: Examples of two core use cases for an AI tutor. **Left:** The *adaptive explanation generation* scenario, showcasing a multi-turn dialogue in electrochemistry. The system must provide a targeted clarification in response to a student’s specific follow-up question regarding standard reduction potentials. **Right:** The *assessment and feedback* scenario, where a student provides a reasoned but incorrect answer to a biology question. The system’s task is to analyze the student’s reasoning, identify the misconception about the Calvin cycle, and provide corrective feedback.

**Use Case 2: Assessment and Feedback.** Students often use LLMs as a means to self-assess their work and get instant feedback. This reflects one of the most promising applications of LLMs: providing real-time evaluation and correction that helps learners iterate quickly. To capture this scenario, TUTORBENCH includes examples where the model is shown a student solution and is asked to assess it. The model is expected to analyze the work, identify any mistakes, classify the nature of each error (e.g., arithmetic, factual, or conceptual), and generate feedback that includes corrections to the student’s solution. An example of this use case is shown in Fig. 2.

**Use Case 3: Active Learning Support.** An essential aspect of high-quality tutoring is to promote active learning, wherein students are encouraged to engage directly with the problem-solving process rather than passively receiving answers. Effective tutors guide students toward the solution through scaffolding, such as hints, analogies, or intermediate steps, which support learning while preserving student agency.

To test this behavior, we present LLMs with incorrect or partially correct student responses and ask them to generate helpful hints that enable the student to take the next step. These hints must strike a careful balance, providing just enough information to resolve confusion without directly giving away the final answer. This task simulates scenarios where a student is stuck midway and requires targeted intervention to proceed. An example of this use case is presented in Fig. 1.

For each use case, models receive the sample along with a system prompt guiding them to act as tutors. Prompts are designed to encourage natural tutoring behaviors without being overly prescriptive (Appendix A.5).

## 2.2 MULTIMODALITY

One key feature of TUTORBENCH is its multimodal design, reflecting the authentic ways students engage with LLMs for tutoring. For each use case illustrated above, we create both text-only and multimodal samples. The images in the dataset consist of hand-written text/diagrams, typed content, or screenshots, mirroring real settings where students often communicate not only through written text but also by sharing images of their work or problem-solving steps. This multimodal component

is essential for evaluating LLMs as effective AI tutors, since it tests their ability to interpret and provide feedback across different forms of student input. An example of a multimodal sample is shown in Fig. 1 where the student solution is in the form of an image of their hand-written work.

### 2.3 LLM JUDGE WITH SAMPLE-SPECIFIC RUBRICS

Given the open-ended nature of the three use cases described above, we adopt a rubric-based evaluation using an LLM-judge based on Claude Sonnet 4 (Anthropic, 2025b) to evaluate model responses. Similar to recent research on rubrics-based model training and evaluation Arora et al. (2025), TUTORBENCH creates separate rubric criteria for each example that help capture fine-grained qualities of a given response. By doing so, we enable efficient and reliable evaluation using an LLM-judge, bypassing a time-consuming human judge for every new model release.

Rubric criteria are designed to be self-contained, mutually exclusive, and collectively comprehensive. This design ensures that they can be applied consistently and without overlap, while also covering the full space of desirable response qualities. Granularity of the criteria makes them particularly well-suited for grading with LLM-based judges. An illustration of the sample-specific rubrics is shown in Fig. 1 (bottom right).

We additionally introduce a weighting scheme to reflect the relative importance of different criteria. Rubrics designated as *critical* are assigned a weight of 5, while *non-critical* rubrics carry a weight of 1. In some cases, *critical* rubrics may be assigned a negative weight of -5 to penalize undesirable behaviors. For instance, in active learning settings, a rubric may require that the final answer not be revealed directly, and a violation of this principle would result in a strong negative score. To quantify overall model performance, we compute a weighted average of the binary scores assigned to each criterion. The final score is normalized to the range [0,1].

We conduct thorough experiments to show the alignment of the LLM-judge vs the human-expert judge. To study the quality of the LLM-judge and compare it with human experts, we collect 3 human ratings per rubric criterion on a subset of 250 samples. Our experiments show that the LLM-judge aligns better with the majority vote (obtained from the 3 human ratings) than the median human expert, and achieves an F1 score of 0.81. Details of experiments are described in Section 4.7.

### 2.4 RUBRIC CRITERIA TAGS

To enable fine-grained analysis of model performance, we annotate each rubric criterion with four tags. Firstly, we annotate each rubric with one of the following *evaluation dimensions* to capture the primary axis along which the model is being judged: instruction following, style and tone, truthfulness, visual reasoning, visual perception, conciseness and relevance, student level calibration, and emotional component. These dimensions provide a structured lens for understanding different facets of tutoring quality. We provide definitions of these dimensions in Appendix A.2. Secondly, we tag the rubric criteria with specific *tutoring skills* being assessed: asking guiding questions, identifying core misconceptions, recognizing correct or incorrect student steps, including examples or analogies, providing alternative solutions, stating definitions or theorems (or other knowledge), or providing step-by-step help. Together, these two tags allow us to analyze not just whether the model performs well, but also what type of tutoring capability is being demonstrated or lacking.

In addition, we annotate each rubric criterion with two complementary tags that enrich the dataset. We annotate each rubric as being *explicit* or *implicit*, to indicate whether the criterion was explicitly requested in the prompt or implicitly inferred from the tutoring context. We also annotate criteria as *objective* or *subjective* to denote whether the criterion judges a response according to an objective standard (e.g., factual accuracy, correctness of a step) or a more subjective one (e.g., appropriateness of tone, empathy). While our primary analyses focus on evaluation dimensions and tutoring skills, we include explicitness and objectivity labels to support further investigations into prompt-response dynamics. We provide details about the distributions of the various tags in Appendix A.4.

## 3 DATA COLLECTION PROCESS

TUTORBENCH consists of high-school-level questions from STEM subjects. The questions and the rubrics to each question were written by experts of the corresponding subject who have a Bachelor’s

or higher degree and have either tutoring or professional experience in the corresponding subject. Each example starts with a question, which is then followed by use-case-specific content. For the adaptive explanation use case, it is followed by (i) a teacher’s explanation of the answer and (ii) a student’s follow-up question asking for clarification on a specific part of that explanation. For the feedback and assessment use case, the initial question is followed by an incorrect solution written in a student persona. For the active learning support use case, the initial question is followed by a partial solution written in a student persona.

Subsequently, the human experts write a list of rubric criteria to evaluate model responses. The rubric criteria are formulated based on a “golden tutoring response” written by them. Each rubric criterion is annotated with several tags as explained in Section 2.4, and assigned a weight to indicate their importance. The weighting system is discussed in Section 2.3.

Five state-of-the-art models, Gemini 2.5 Pro, GPT-o3, Claude 3.7 Sonnet , DeepSeek-R1, Llama 4 Maverick are then prompted with the full example, along with use-case-specific instructions. The obtained responses are then graded against the rubric criteria. We keep only those samples for which 3 out of the 5 models achieve less than 50% score (weighted) across the rubric criteria.

The resulting dataset comprises 1,490 examples distributed across six STEM subjects: physics, chemistry, biology, calculus, computer science, and statistics, with a relatively balanced representation. The tasks are evenly split across the three use cases, with each use case also including a substantial portion of multimodal samples. In total, 828 examples are multimodal, combining text with images of student work, while the remainder are text-only. The distribution of samples across subjects, use cases, and modality is shown in Fig. 7 in Section A.1. Each example in the dataset is associated with a separate set of rubric criteria, ranging from 3 to 39 per example, providing fine-grained evaluation criteria. This results in a total of 15,220 rubrics across the whole dataset.

## 4 EVALUATION AND ANALYSIS

In the following subsections, we present the performance of a select set of state-of-the-art LLMs on TUTORBENCH in terms of their overall performance, and more fine-grained performance along use cases and other dimensions.

### 4.1 OVERALL MODEL PERFORMANCE

Model performance on TUTORBENCH is evaluated using the weighted average rubric rating  $\text{ARR}_w$ , defined on the  $j$ -th example as

$$\text{ARR}_w^j = \frac{\sum_{i=1}^{N_j} w_i^j \cdot \mathbb{1}_{r_i^j}}{\sum_{i=1}^{N_j} w_i^j \cdot \mathbb{1}_{w_i^j > 0}}, \quad (1)$$

where  $N_j$  is the number of rubric criteria in the  $j$ -th example,  $w_i^j \in \{-5, 1, 5\}$  is the weight of the  $i$ -th rubric for the  $j$ -th test example, and  $r_i^j \in \{0, 1\}$  is the fail/pass rating of the model response on the  $i$ -th rubric of the test example  $j$ . The final score for a model is then the average  $\text{ARR}_w$  across all the examples.

The evaluation results are presented in Table 1. We observe that Gemini 2.5 Pro (Google DeepMind, 2025) and GPT-5 (GPT-5, 2025) obtain the best overall performance, with very similar scores on both text-only and multimodal tests. Interestingly, OpenAI and Gemini models outperform all the Claude model series, including the latest Claude Opus 4.1 (Thinking). Furthermore, none of the models surpass 56% overall performance, highlighting the complex nature of TUTORBENCH. Finally, it is noteworthy that the recently released gpt-oss-120b model performs close to the best models on the text-only subset. We are unable to measure its performance on the multimodal subset, since it does not support image input.

### 4.2 PERFORMANCE BY USE CASE

Performance breakdown by use case is shown in Fig. 3. On average, models achieve a score of 47.16% on adaptive explanation generation, 51.56% on assessment and feedback, and 54.07% on

Rank	Model	Text-Only (%)	Multimodal (%)	Overall (%)	CI
1	Gemini 2.5 Pro	<b>57.05</b>	<b>54.53</b>	<b>55.65</b>	$\pm 1.11$
2	GPT-5	<b>57.03</b>	53.97	55.33	$\pm 1.02$
3	o3 Pro	56.07	53.45	54.62	$\pm 1.02$
4	o3 Medium Effort	54.11	51.68	52.76	$\pm 1.00$
5	o3 High Effort	52.91	51.43	52.09	$\pm 1.01$
6	Claude Opus 4.1 (Thinking)	51.65	50.08	50.78	$\pm 1.05$
7	Claude Opus 4 (Thinking)	50.40	49.14	49.71	$\pm 1.02$
8	Claude Opus 4.1	49.51	45.72	47.40	$\pm 1.06$
9	Claude 3.7 Sonnet (Thinking)	45.67	47.07	46.45	$\pm 1.03$
10	Claude Opus 4	47.79	43.59	45.46	$\pm 1.06$
11	Llama 4 Maverick	39.54	40.73	40.20	$\pm 1.00$
12	GPT-4o	39.10	33.74	36.12	$\pm 0.96$
	gpt-oss-120b	56.01	N/A	N/A	$\pm 1.49$
	gpt-oss-20b	49.01	N/A	N/A	$\pm 1.53$
	DeepSeek-R1	48.38	N/A	N/A	$\pm 1.50$

Table 1: Evaluation of state-of-the-art LLMs on TUTORBENCH. Gemini 2.5 Pro achieves the highest score (55.65%), highlighting the complex nature of the samples in TUTORBENCH and the need for further advancement in models for tutoring applications. We report scores on the text-only samples for gpt-oss-120b, gpt-oss-20b, and DeepSeek-R1 because they do not support multimodal input and were therefore evaluated only on the text-only subset.

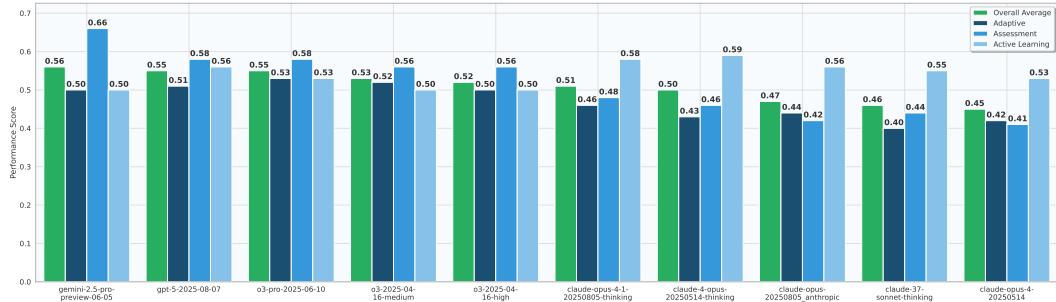


Figure 3: Model performance across three use cases: **Adaptive**, **Assessment**, and **Active Learning**. We observe a distinct difference between the performance of the Claude family of models compared to the other models, with the Claude models performing significantly better in providing active learning support, but still lagging behind other models overall.

active learning support. We also observe that the Claude family of models performs much better in the active learning support use case compared to the other models, although their overall performance is poorer.

#### 4.3 PERFORMANCE ALONG EVALUATION DIMENSIONS

As explained in Section 2.4, we annotate each rubric criterion with one **evaluation dimensions**. This allows us to aggregate the pass/fail ratings on subsets with each of the tags. We present the mean pass/fail rate along these dimensions for the top 10 models (according to Table 1) in Fig. 4. While performance along these dimensions roughly follows the same trend as the overall performance, Gemini 2.5 Pro shows a much better performance than other models in recognizing student emotions such as confusion, frustration, curiosity, and generates responses with the right tone and style (e.g., by using headings, bullets, and LaTeX). On the other hand, GPT-5 and o3 Pro perform best on factual correctness (truthfulness), student-level calibration, and instruction following.

#### 4.4 PERFORMANCE ALONG DIFFICULTY LEVEL

The study categorizes evaluation samples using Bloom’s taxonomy, which organizes cognitive tasks from basic skills like remembering and understanding to advanced ones such as analyzing, evalu-

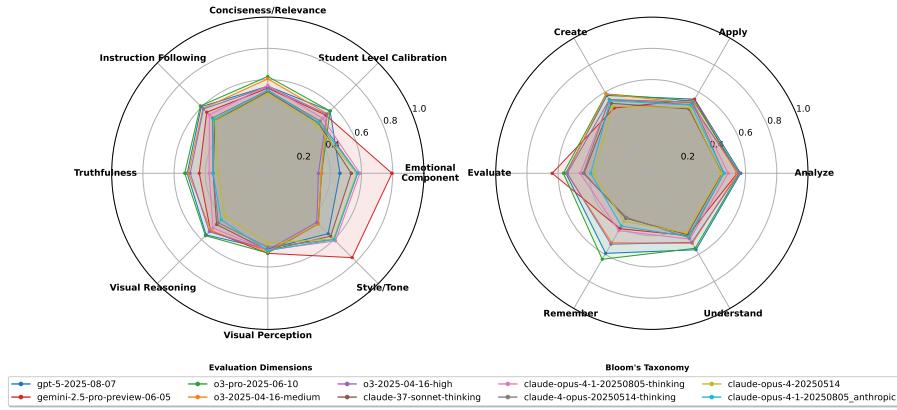


Figure 4: Model performance breakdown along evaluation dimensions and Bloom’s taxonomy categories. While the top-performing models GPT-5 and Gemini 2.5 Pro are close overall, their performance differs widely when measured along the above dimensions.

ating, and creating, in order to better assess model performance across different levels of cognitive demand. Three state-of-the-art LLMs (Gemini 2.5 Pro, Claude Sonnet 4, and GPT-5) were used to annotate each sample, with a category label assigned if at least two models agreed, covering over 97% of the samples. Final scores for each category were calculated by averaging rubric-based scores across all annotated samples. Results, shown in Figure 4, reveal that performance does not align with the taxonomy’s order of difficulty. For example, Gemini 2.5 Pro scores lower on “remember” than on “evaluate”, suggesting that while models may demonstrate advanced reasoning, they can struggle with recalling and explicitly stating content, which is crucial for tutoring tasks that require effective context-driven communication.

#### 4.5 PERFORMANCE ALONG TUTORING SKILLS

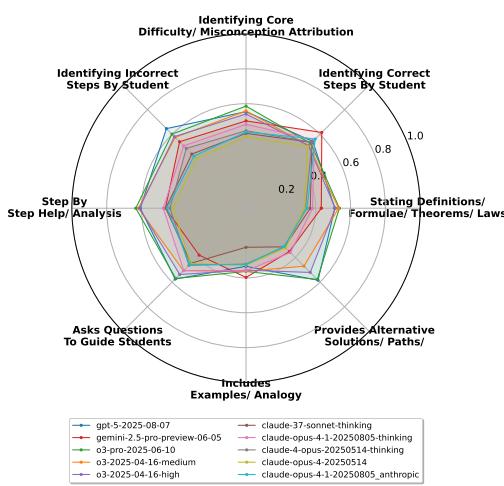


Figure 5: Model performance breakdown along tutoring skills: models struggle to include alternative solutions, examples, and analogies in their responses. However, they perform relatively better in identifying mistakes, correct steps, and core misconceptions.

Figure 5 illustrates the performance of various models across eight distinct tutoring skills. The skills are: Identifying Core Difficulty/Misconception Attribution, Identifying Incorrect Steps By Student, Identifying Correct Steps By Student, Stating Definitions/Formulas/Theorems/Laws, Provides Alternative Solutions/Paths/, Includes Examples/Analogy, Asks Questions To Guide Students, and Step By Step Help/Analysis. The chart shows that models generally perform well in identifying correct steps and student errors but struggle with providing alternative solutions and examples.

Tutoring is a complex task requiring several nuanced skills, such as calibrating to the student, recalling and applying knowledge, or breaking down problems into simpler steps. To evaluate models along with such finer skills, we classify the rubric criteria written by human-experts into 8 high-level skills: identifying core difficulty, identifying students’ correct steps, identifying students’ incorrect steps, recalling and stating knowledge, providing alternative solutions, including examples, asking questions to guide students, and providing step-by-step help. These skills were identified by manually inspecting 600 rubric criteria and by using an LLM to tag all the rubric criteria with one of the skills, if applicable.

We present the average pass/fail rate of LLM responses on these tags in Figure 5. Overall, we observe that models perform well in identifying correct/incorrect steps by students (average scores of 53.7% and 53%), but they struggle to include alternate solutions, and examples/analogies in their responses (achieving average scores of 41.9% and 32.8%). Model performance varies across skills, with Gemini 2.5 Pro excelling at identifying correct/ incorrect steps and providing examples, while GPT-5 is stronger at spotting errors and misconceptions.

## 4.6 AN EVALUATION OF GPT-5 STUDY MODE

Along with the models in Table 1, we also evaluated OpenAI’s recently released *study mode* OpenAI (2025b). According to their website, study mode is “powered by custom system instructions written in collaboration with teachers, scientists, and pedagogy experts” to promote deeper learning. Since it is only available via the web interface, we collected responses from the website.

Unlike API models, study mode does not allow direct system prompts and instead engages users in back-and-forth dialogue calibrated to their skill level. This makes apples-to-apples comparison with TUTORBENCH difficult, since TUTORBENCH evaluates single, pre-formulated responses. For this reason, we exclude study mode from the main leaderboard and report results separately.

On TUTORBENCH, study mode scores  $46.94 \pm 1.06\%$  overall (48.19% text-only, 45.93% multi-modal), indicating poor performance against rubric criteria. Comparing with API-based GPT-5, we note two key observations: (1) study mode often provides only partial, intermediate responses, typically ending with a counter-question (e.g., “Before we move on, do you want me to walk you through …?”). These incomplete answers naturally score lower. (2) Even when complete, study mode responses perform worse than GPT-5’s, as they are more concise and address fewer rubric criteria (examples in Appendix A.3).

Thus, study mode’s lower scores stem both from TUTORBENCH ’s unsuitability for partial responses and from study mode’s tendency to neglect several rubric criteria even in complete answers.

## 4.7 QUALITY OF LLM-JUDGE USED IN AUTOMATED EVALUATIONS

The rubric criteria in TUTORBENCH cover several criteria that together reflect an ideal response. To achieve a scalable evaluation, we rely on an LLM-judge to grade model responses against the criteria. We experimented with multiple state-of-the-art models and chose Claude Sonnet 4 as our judge model, as it achieved the best alignment with human ratings.

To compare the LLM-judge with a human expert, we first collect 3 separate human-expert ratings across 250 examples, and a total of 2475 rubric criteria. To capture the inter-human ratings agreement, we compute pairwise agreements between the 3 ratings for all rubrics and compute the average agreement rate. The mean inter-human agreement across the 69 experts was 0.75, while agreement between the LLM-judge and human ratings was 0.78. The distribution of agreements per individual human expert who contributed to our annotations is shown in Fig. 6, along with the LLM-judge agreement.

We also study the alignment of the LLM-judge with the majority vote on rubrics. Firstly, we filter out the *critical* rubric criteria (with weights +5 or -5) from the 250 examples, resulting in a total of 1900 criteria with 3 ratings each. We then use the majority vote (pass/fail) as the label for each rubric. The LLM judge achieves an F1-score of 0.82. The highest F1-score achieved by a single human-expert was 0.91, computed across a set of 321 criteria that were rated by them. This demonstrates that the LLM-judge has a strong performance and aligns well with human-expert ratings.

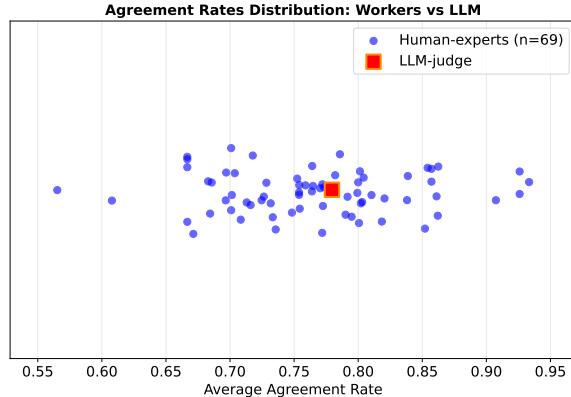


Figure 6: We measure the quality of the LLM-judge used in our evaluations by comparing its rating on model responses across 250 samples from TUTORBENCH with the majority vote obtained using 3 human-expert ratings. We observe that the LLM-judge ranks better than the median human-expert, thus demonstrating a strong alignment of the LLM-judge with human ratings.

## 5 RELATED WORK

**Evaluation Frameworks for AI-Tutor Models.** Evaluating tutoring quality has emerged as an important area of study. Maurya et al. (2025) propose MRBench, an evaluation benchmark with 192 math tutoring conversations. They propose a uniform taxonomy for evaluating tutoring capabilities based on learning principles. In a similar effort to evaluate the impact of using LLM-based tutors, Lyu et al. (2024) conduct a field study of 50 students by using an LLM-based coding assistant. While the students who used the assistant achieved statistically significant improvement in their performance, they expressed concerns about the limited role the assistant played in developing critical thinking skills. Similarly, Liu et al. (2024) document the integration of AI-based tools into Harvard’s CS50 course. These classroom studies reveal how AI tutors can complement human instructors at scale but often remain limited to narrow domains (e.g., math word problems, computer science) and lack multimodal or rubric-based evaluation.

**LLMs as Tutors.** Recent research has also explored designing improved tutors using LLMs. Pal Chowdhury et al. (2024) develop a structured, rule-based tutoring framework with LLMs. Their system embeds guardrails and predefined pedagogical strategies, allowing LLMs to generate tutor responses while maintaining instructional control. Compared to free-form GPT-4 tutors, this hybrid approach reduced redundant or unhelpful dialog turns and better adhered to pedagogical principles. Similarly, Wang et al. (2024) propose Bridge, which uses cognitive task analysis to capture expert teachers’ decision-making processes, such as diagnosing student errors and selecting remediation strategies, before feeding them into LLMs.

In contrast to these efforts, our dataset targets multimodal tutoring across six high school STEM subjects, moving beyond single-domain math word problems. Our work also introduces explicit rubric-based evaluation criteria tied to each sample, enabling efficient and scalable evaluation of multiple models. Our dataset is also substantially larger and covers a diverse set of use cases, topics, and evaluation dimensions. Overall, TUTORBENCH provides a broader testbed for measuring and advancing LLMs’ tutoring capabilities.

## 6 LIMITATIONS OF TUTORBENCH

While our dataset provides a structured foundation for evaluating LLM tutoring, it has important limitations. Its scope of instructional use cases is constrained: we focus on three representative tutoring scenarios but omit other valuable tasks such as generating practice problems, designing exercises, or introducing new concepts. The dataset also evaluates only final responses to pre-formulated conversations, limiting assessment of adaptability in dynamic, multi-turn exchanges.

In addition, while learning is multimodal, our dataset only incorporates images of students’ work as input and does not test visual content generation. Our STEM focus provides depth in reasoning and problem-solving but excludes humanities domains that rely on narrative and interpretive skills. Finally, we isolate model outputs from the broader UI context, though real-world tutoring effectiveness depends on presentation, interactivity, and learner engagement. Future benchmarks could build on this dataset with scaffolding to address these interaction effects.

## 7 CONCLUSION

We introduce TUTORBENCH, a dataset and evaluation benchmark to assess tutoring capabilities of LLMs. We focus on three specific user scenarios and study how models perform when asked to (i) adapt their explanations to a student with a given background, (ii) provide an assessment of students’ work, and (iii) provide hints to support active learning. TUTORBENCH is a rubrics-based evaluation and each sample in the dataset is accompanied by specific rubric criteria. These rubric criteria are weighted based on importance and are also tagged with several attributes that allow for finer analysis of model behavior. Another important aspect of TUTORBENCH is the support for multimodal input: the dataset includes samples with images containing students’ work to reflect real-world usage. We also provide a comprehensive evaluation of a select set of state-of-the-art LLMs and identify growth areas. Overall, TUTORBENCH serves as one of the first datasets to comprehensively evaluate LLMs on tutoring-based applications.

## REFERENCES

- Tawfiq Ammari, Meilun Chen, S. M. Mehedi Zaman, and Kiran Garimella. How students (really) use chatgpt: Uncovering experiences among undergraduate students. *arXiv preprint arXiv:2505.24126*, 2025. doi: 10.48550/arXiv.2505.24126.
- Anthropic. Claude 3.7 sonnet and claude code, February 2025a. URL <https://www.anthropic.com/news/clause-3-7-sonnet>. Product announcement.
- Anthropic. Claude sonnet 4: Hybrid reasoning model from anthropic. <https://www.anthropic.com/clause/sonnet>, May 2025b. Also detailed in the Claude 4 announcement (May 22 2025) and Wikipedia entry (accessed 2025-09-06).
- Rahul K Arora, Jason Wei, Rebecca Soskin Hicks, Preston Bowman, Joaquin Quiñonero-Candela, Foivos Tsimpourlas, Michael Sharman, Meghan Shah, Andrea Vallone, Alex Beutel, et al. Health-bench: Evaluating large language models towards improved human health. *arXiv preprint arXiv:2505.08775*, 2025.
- Hamsa Bastani, Osbert Bastani, Alp Sungu, Haosen Ge, Özge Kabakçı, and Rei Mariman. Generative ai without guardrails can harm learning: Evidence from high school mathematics. *Proceedings of the National Academy of Sciences*, 122(26):e2422633122, 2025. doi: 10.1073/pnas.2422633122. URL <https://www.pnas.org/doi/abs/10.1073/pnas.2422633122>.
- DeepSeek. Deepseek-r1 release, January 2025. URL <https://api-docs.deepseek.com/news/news250120>. Model announcement and licensing details.
- Sean Geraghty and Nash Goldstein. 11th grader takes an ai tutoring deep dive: And a human tutoring expert takes notes. Education Next (blog post), November 22 2024. URL <https://www.educationnext.org/11th-grader-takes-an-ai-tutoring-deep-dive/>. Last updated November 22, 2024.
- Elliot Glazer, Ege Erdil, Tamay Besiroglu, Diego Chicharro, Evan Chen, Alex Gunning, Caroline Falkman Olsson, Jean-Stanislas Denain, Anson Ho, Emily de Oliveira Santos, Olli Järvinieemi, Matthew Barnett, Robert Sandler, Matej Vrzala, Jaime Sevilla, Qiuyu Ren, Elizabeth Pratt, Lionel Levine, Grant Barkley, Natalie Stewart, Bogdan Grechuk, Tetiana Grechuk, Shreepranav Varma Enugandla, and Mark Wildon. Frontiermath: A benchmark for evaluating advanced mathematical reasoning in ai. *arXiv preprint arXiv:2411.04872*, November 2024. doi: 10.48550/arXiv.2411.04872.
- Google. Google gemini: Free pro plan for students. <https://gemini.google/students/>, 2025. Accessed September 12, 2025. Offer ends October 6, 2025.
- Google DeepMind. Gemini 2.5: Our most intelligent ai model, March 2025. URL <https://blog.google/technology/google-deepmind/gemini-model-thinking-updates-march-2025/>. Announcing availability of Gemini 2.5 Pro.
- OpenAI GPT-5. GPT-5 System Card. Online; August 7, 2025, 2025. URL <https://openai.com/index/gpt-5-system-card/>. Accessed: 2025-08-27.
- Adit Gupta, Jennifer Reddig, Tommaso Calo, Daniel Weitekamp, and Christopher J. MacLellan. Beyond final answers: Evaluating large language models for math tutoring. *arXiv preprint arXiv:2503.16460*, February 2025. doi: 10.48550/arXiv.2503.16460.
- Kunal Handa, Drew Bent, Alex Tamkin, Miles McCain, Esin Durmus, Michael Stern, Mike Schiraldi, Saffron Huang, Stuart Ritchie, Steven Syverud, Kamya Jagadish, Margaret Vo, Matt Bell, and Deep Ganguli. Anthropic education report: How university students use claude, 2025. URL <https://www.anthropic.com/news/anthropic-education-report-how-university-students-use-claude>.

Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, Junhao Shen, Jinyi Hu, Xu Han, Yujie Huang, Yuxiang Zhang, Jie Liu, Lei Qi, Zhiyuan Liu, and Maosong Sun. Olympiadbench: A challenging benchmark for promoting agi with olympiad-level bilingual multimodal scientific problems. *arXiv preprint arXiv:2402.14008*, February 2024. doi: 10.48550/arXiv.2402.14008.

Rongxin Liu, Carter Zenke, Charlie Liu, Andrew Holmes, Patrick Thornton, and David J. Malan. Teaching cs50 with ai: Leveraging generative artificial intelligence in computer science education. In *Proceedings of the 55th ACM Technical Symposium on Computer Science Education V. 1*, SIGCSE 2024, pp. 750–756, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400704239. doi: 10.1145/3626252.3630938. URL <https://doi.org/10.1145/3626252.3630938>.

Wenhan Lyu, Yimeng Wang, Tingting (Rachel) Chung, Yifan Sun, and Yixuan Zhang. Evaluating the effectiveness of llms in introductory computer science education: A semester-long field study. In *Proceedings of the Eleventh ACM Conference on Learning @ Scale, L@S '24*, pp. 63–74, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400706332. doi: 10.1145/3657604.3662036. URL <https://doi.org/10.1145/3657604.3662036>.

Kaushal Kumar Maurya, Kv Aditya Srivatsa, Ksenia Petukhova, and Ekaterina Kochmar. Unifying AI tutor evaluation: An evaluation taxonomy for pedagogical ability assessment of LLM-powered AI tutors. In Luis Chiruzzo, Alan Ritter, and Lu Wang (eds.), *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 1234–1251, Albuquerque, New Mexico, April 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.nacl-long.57. URL <https://aclanthology.org/2025.nacl-long.57/>.

Meta AI. The llama 4 herd: The beginning of a new era of natively multimodal intelligence, April 2025. URL <https://ai.meta.com/blog/llama-4-multimodal-intelligence/>. Announcing Llama 4 Scout and Llama 4 Maverick.

OpenAI. Introducing openai o3 and o4-mini, April 2025a. URL <https://openai.com/index/introducing-o3-and-o4-mini/>. Model announcement; later updates noted on page.

OpenAI. Introducing study mode. <https://openai.com/index/chatgpt-study-mode/>, July 2025b. Accessed September 12, 2025. Feature launched July 29, 2025.

Sankalan Pal Chowdhury, Vilém Zouhar, and Mrinmaya Sachan. Autotutor meets large language models: A language model tutor with rich pedagogy and guardrails. In *Proceedings of the Eleventh ACM Conference on Learning @ Scale (L@S '24)*, pp. 5–15, Atlanta, GA, USA, 2024. Association for Computing Machinery. doi: 10.1145/3657604.3662041.

Long Phan, Alice Gatti, Ziwen Han, Nathaniel Li, Josephina Hu, Hugh Zhang, et al. Humanity's last exam. *arXiv preprint arXiv:2501.14249*, January 2025. doi: 10.48550/arXiv.2501.14249.

Alexander Scarlatos, Naiming Liu, Jaewook Lee, Richard Baraniuk, and Andrew Lan. Training llm-based tutors to improve student learning outcomes in dialogues. *arXiv preprint arXiv:2503.06424v1*, March 2025. doi: 10.48550/arXiv.2503.06424.

Alessandro Vanzo, Sankalan Pal Chowdhury, and Mrinmaya Sachan. GPT-4 as a homework tutor can improve student engagement and learning outcomes. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (eds.), *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 31119–31136, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-251-0. doi: 10.18653/v1/2025.acl-long.1502. URL [https://aclanthology.org/2025.acl-long.1502/](https://aclanthology.org/2025.acl-long.1502).

Rose Wang, Qingyang Zhang, Carly Robinson, Susanna Loeb, and Dorottya Demszky. Bridging the novice-expert gap via models of decision-making: A case study on remediating math

mistakes. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, Mexico City, Mexico, June 2024. Association for Computational Linguistics. URL <https://aclanthology.org/2024.naacl-long.120/>.

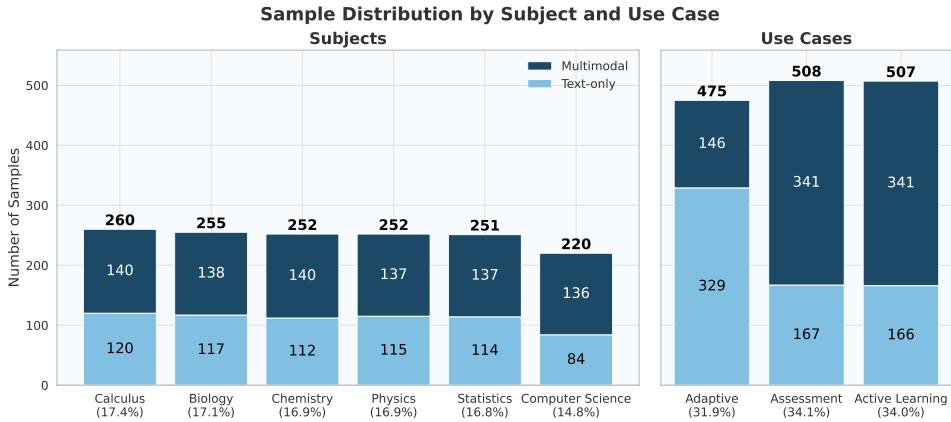


Figure 7: Distribution of samples by subject (left) and use case (right). The stacked bars indicate the number of Text-only (light blue) and Multimodal (dark blue) samples in each category. The distribution across the six subjects is relatively balanced, while the counts for the use cases vary more significantly. The labels ‘Adaptive’, ‘Assessment’, and ‘Active Learning’ correspond to adaptive explanation generation, assessment and feedback, and active learning support, respectively.

## A APPENDIX

### A.1 DISTRIBUTION OF SAMPLES OVER SUBJECTS AND USE CASES

### A.2 RUBRIC CRITERIA TAG DEFINITIONS

As explained in Section 2.4, we annotate each rubric criterion with several tags to enable fine-grained analysis of model behavior. In Table 2, we provide definitions of the *evaluation dimensions* tag.

Tag	Definition
Instruction following	Criterion checks for adherence to system instructions and prompt instructions
Truthfulness	Criterion checks for factual accuracy, capturing the “precision” of the model, and does not penalize missed content, incorrect formatting
Conciseness and Relevance	Criterion checks for how direct, on-topic, and efficient the content is. Examples are criteria that check for what should not be in the response, unnecessary formatting
Style and Tone	Criterion checks for clarity, fluency, and appropriateness of tone
Visual Perception [Multimodal Only]	Criterion checks whether the model correctly identifies the content from the image required to solve the problem.
Visual Reasoning [Multimodal Only]	Criterion checks whether the model correctly reasons based on the provided image. If the reasoning does not depend on the image, this tag is not used.
Student-level Calibration	Criterion checks if the model response accounts for the explicit/ implicit knowledge level of the student.
Emotional Component Calibration	Criterion checks if the model pays attention to a student’s enthusiasm, frustrations, overwhelm, confusion, and other emotional cues mentioned in the prompt.

Table 2: Definitions of the evaluation dimensions used to tag rubric criteria.

### A.3 GPT-5 STUDY MODE RESPONSE EXAMPLES

In Section 4.6, we provided details about why the study mode responses from GPT-5 collected from the website rank lower than GPT-5 itself. Here, we provide illustrative examples of cases where either the study mode response is indeed poorer or is truncated due to the model seeking user input

Question: What is responsible for the seasonal change in the coat color of Arctic foxes?

In winter, Arctic foxes usually have a white coat. That's because lower temperatures in winter denature the pigment molecules in the fox's coat, causing it to be lighter in color. In summer, when the temperature rises again, |

Study mode (web)

GPT-5 (API)

Think about how temperature affects pigment production in the fox's fur — what happens to those pigment molecules in the warmer months compared to the colder months? That contrast will help you explain the seasonal color change.

Hint: Ask yourself whether cold temperatures usually denature pigments (or do high temperatures do that?). Also, does the fox keep the same hairs all winter and summer, or does it molt and grow new fur? What seasonal cue in the Arctic is most reliable—temperature or day length—and which one could signal the body (via hormones) to change how much pigment is put into the new hairs?

	Study mode (web)	GPT-5 (API)
The response must acknowledge the student's current work by referring to their claim that lower temperatures in winter denature the pigment molecules in the fox's coat, causing it to be lighter in color.	0	0
The response must identify that the student is stuck due to a conceptual misunderstanding of what actually causes denaturation (i.e., high rather than low temperatures and extreme pH).	0	1
The response must hint that the seasonal coat-color changes observed in the Arctic fox cannot be caused by denaturation by encouraging the student to consider whether denaturation could occur at sub-zero arctic temperatures.	0	1
The response must hint that the seasonal coat-color changes observed in the Arctic fox cannot be caused by denaturation by encouraging the student to consider whether hair follicles whose pigment molecules are irreversibly denatured would ever be able to produce dark hairs again in the summer.	0	0
The response must provide a hint promoting the student to consider that denaturation is irreversible and therefore unlikely to explain the seasonal coat-color change.	0	0
The response must hint that a hormone affected by day length changes plays a role in regulating melanin production in Arctic foxes.	0	1
The response must hint that changes in day length can biologically influence fur color through melanin production.	0	1
The model must be able to reason that the student's response that the Arctic fox's pigment molecules are denatured in the winter is incorrect because denaturation cannot occur at the extreme low temperatures in which the Arctic fox lives.	0	1
The response should mention that the Arctic fox grows a new coat each season, which may help explain reversible changes in coat color.	0	1
The response must not state that the correct final answer is that the seasonal coat-color changes observed in the Arctic fox are due to photoperiodic melatonin regulation of melanin.	1	1

Figure 8: An example comparing study mode and GPT-5 on the use case of *active learning support*. We can observe that the rubrics require the response to correct a misconception in the student regarding the relationship between temperature and fur color in arctic foxes and guide the student to consider other factors (such as daylight duration and hormonal changes) affecting the fur color during winter. As observed, the response from GPT-5 (API) addresses the rubric criteria directly, while the response from study mode is brief and does not refer to any other factors at all.

before proceeding. In the latter case, the model response is incomplete (probably by design) and hence scores poorly against the rubric criteria. We present three examples in Figs. 8 to 10 where their responses are shown side-by-side along with the corresponding rubric criteria. These examples highlight some of the shortcomings of study mode: the responses are sometimes terse, or fail to get into the details of the problem, or sometimes expect a response from the student before proceeding. By grounding the comparison in rubric-based criteria, we provide concrete evidence for why study mode ranks lower in Table 1.

A right prism has a base that is an isosceles trapezoid with legs that are \$5\$ inches in length and with one parallel side that is \$6\$ inches longer than the other parallel side. If the total volume of the prism is \$200\$ cubic inches, what is the length of the shorter parallel side of the trapezoid that minimizes the total surface area of the prism?

Initial solution (content omitted here for brevity)

I'm really confused. I don't understand the surface area calculation. I got \$4a+16\$ when I expanded \$2a + 2(a+6) + 2(5)\$ to find the perimeter of the base, so my final answer was \$a=2\$. What did I do wrong?

**Study mode (web)**

Great question—this snafu is all about “what gets doubled” in the surface area formula of a prism.

For a right prism:

“Surface area” =  $2 \times (\text{base area}) + (\text{perimeter of ONE base}) \times (\text{height of prism})$ . Only the “base area” is doubled (top + bottom). The “perimeter” is “not” doubled.

In your work you used something like  $2a + 2(a+6) + 2(5)$ . Those extra 2’s are the issue—you accidentally doubled each side as if there were two perimeters.

Quick check together: for the isosceles trapezoid base with sides  $a, a+6, 5, 5$ , “what is the perimeter of a single base” (no doubling)? Just add them once:  $\text{Text}(\text{Perimeter}) = 12$ .

(Reply with that number, and we'll plug it into the correct surface area expression step by step.)

**GPT-5 (API)**

Great question—you’re right to focus on the surface area setup. Here’s where the mix-ups happened and how to fix them.

1) Which perimeter to use  
For a right prism, the total surface area is:  
 $SA = 2B + PH$   
-  $B$  = area of one base  
-  $P$  = perimeter of one base (not doubled)  
-  $H$  = prism height

So you should use  $P = a + (a+6) + 5 + 5 = 2a + 16$ . Writing  $2a + 2(a+6) + 2(5)$  actually gives  $2P$ , which double-counts the perimeter.

2) Small algebra slip  
Even if you were computing  $2P$ , expanding  $2(a+6)$  should give  $2a + 12$ , not  $2a + 6$ . So  $2a + 2(a+6) + 2(5) = 2a + (2a + 12) + 10 = 4a + 22$ , not  $4a + 16$ .

3) Correct setup and result  
- Height of trapezoid base:  $ht = 4$  (from  $5^2 = ht^2 + 3^2$ ).  
- Base area  $B = \frac{1}{2}(a + a + 6)4 = 4a + 12$ .  
- Prism height from volume  $200 = B H$ :  $H = 200/(4a + 12) = 50/(a + 3)$ .  
- Perimeter:  $P = 2a + 16$ .

Surface area:  
 $SA = 2B + PH = 2(4a + 12) + (2a + 16)50/(a + 3)$   
 $= 8a + 24 + 1000/(a + 3)$ .

Differentiate and set to zero:  
 $d(SA)/da = 8 - 500/(a + 3)^2 = 0$   
 $\Rightarrow (a + 3)^2 / 500 = 62.5$   
 $\Rightarrow a + 3 = \sqrt{62.5} = 7.91$  inches.

Quick check: with  $a = 2$  (your answer),  $SA = 240$ ; with  $a = 4.91$ ,  $SA = 227$ , which is smaller. So the minimum occurs at  $a = 4.91$  inches.

Does the  $SA = 2B + PH$  formula and the perimeter  $P = 2a + 16$  make sense now?

	Study mode (web)	GPT-5 (API)
The response must acknowledge the student's explicitly shown feeling of confusion.	1	1
The response must identify that the student's missing background knowledge was in how to use the distributive property when simplifying the expression $2a + 2(a+6) + 2(5)$ .	0	0
The response must identify that the student made an error when simplifying the expression $2a + 2(a+6) + 2(5)$ and getting $4a + 16$ .	0	1
The response must explain that the distributive property must be used when multiplying a number by a sum in parentheses, so the number must be multiplied by both terms in parentheses.	0	0
The response must correct the student's error by giving the correct expansion for $2a + 2(a+6) + 2(5)$ , which is $4a + 22$ .	0	1
The response must identify that the student's final answer of $2$ is incorrect.	0	1
The response must explain that when calculating the surface area, the sides ( $a, a+6, 5, 5$ ) are added together to get the perimeter of the base, leading to a correct surface area equation of $2(4a+12) + (2a+16)\frac{50}{a+3}$ with derivative $8 - \frac{500}{(a+3)^2} = 0$ .	0	1
The response must give a correct final answer of $\sqrt{62.5} - 3$ or an approximation of approximately $4.91$ .	0	1

Figure 9: An example comparing study mode and GPT-5 on the use case of *adaptive explanation generation*. We can observe that the GPT-5 response clearly identifies that the student is confused about the formula of the perimeter and offers clarification, unlike study mode. Further, study mode does not provide the correct answer, but expects the student to respond before proceeding further.

#### A.4 DISTRIBUTION OF RUBRIC CRITERIA TAGS

In this section, we present the distribution of the four tags that each rubric criterion was annotated with when applicable. In Fig. 11, we show the distribution of the objectivity and explicitness tags. We observe that the majority of the rubric criteria are both objective and explicit. In Fig. 12a, we show the distribution of the evaluation dimensions, and in Fig. 12b, we show the distribution of tutoring skills across the rubrics.

#### A.5 SYSTEM PROMPTS

To generate model responses for evaluation, we use system prompts to lightly guide the LLMs to assume the role of a tutor. We refrain from using overly prescriptive system prompts to measure the



**Study mode (web)**

```

Great! Let's go through the solution step-by-step to assess it.

### What the solution does well:
- Correctly identifies that the residual is "actual value - predicted value".
- Calculates the predicted weight by plugging the chest circumference into the regression line.
- Shows the substitution and arithmetic clearly.
- Correctly computes the residual as approximately 10 kg.
- Interprets the residual meaningfully: the model underestimates the elk's weight by about 10 kg.

### Minor points or improvements:
- The solution rounded the multiplication $3.7455 \times 145.9$ to 545.0, which is a reasonable approximation, but you could check the exact value for precision.
- The last comment about dividing by predicted value (to get a relative error) is not necessary for the residual here. Residuals in regression are simply the difference (actual - predicted), so leaving it as is is correct.
- The answer clearly shows the reasoning and method, which is excellent.

### Summary assessment:
**This is a solid and accurate solution.** The student correctly computed the predicted weight, calculated the residual as actual minus predicted, and interpreted it properly. The minor rounding does not significantly affect the correctness.

If you'd like, I can help you refine or explain any part of this solution. How confident do you feel about residuals in regression? Would you like a quick recap?

```

**GPT-5 (API)**

Following is the equation of the least-squares regression line relating chest circumference and weight for male tule elk: predicted weight =  $-350.3 + 3.7455(\text{chest circumference})$ . The weight of one male tule elk with a chest circumference of 145.9 cm is 204.3 kg. Calculate the residual for this male elk, making sure to show your work along the way.

Hmm ... to get the residual I first need the predicted weight, so I will plug the elks chest size into the regression line:  $w^* = -350.3 + 3.7455 * (145.9)$ . I rounded  $3.7455 * 145.9$  to about 545.0, so  $w^* \approx -350.3 + 545.0 = 194.7$  kg.

Residuals measure how far the model misses, and my notes say it is actual - predicted (so positive means the elk is heavier than the line guessed). Therefore residual =  $204.3 - 194.7 \approx 10$  kg.

I think that means the model under-estimated this elks weight by roughly 10 kilograms. (I am not completely certain if I had to divide by the predicted value to make it a relative error, but I will leave it as the difference for now.)

Assessment: Partially correct method, but numerical error and some confusion.

What you did well:

- You used the right idea: compute the predicted weight from the regression line, then residual = actual - predicted.
- You interpreted the sign correctly: a positive residual means the elk is heavier than the model predicted.

Mistakes and corrections:

- Your multiplication/rounding was too rough. Compute the predicted weight more accurately:
  - $y\text{-hat} = -350.3 + 3.7455(145.9)$
  - $3.7455 \times 145.9 = 546.46845$
  - $y\text{-hat} = -350.3 + 546.46845 = 196.16845 \text{ kg} (\approx 196.17 \text{ kg})$
- Residual = actual - predicted =  $204.3 - 196.16845 = -8.13155 \text{ kg} (\approx -8.13 \text{ kg})$ 
  - You reported about 10 kg, which is off by nearly 2 kg—larger than reasonable rounding error given a calculator should be used here.**
- You don't divide by the predicted value for a residual. That would be a relative (percent) error, not the residual. A standardized residual divides by the residual's standard deviation, but that's not what's asked.

Final: The correct residual is approximately  $-8.13 \text{ kg}$ , indicating the model underestimated this elk's weight by about  $8.1 \text{ kg}$ .

	<b>Study mode (web)</b>	<b>GPT-5 (API)</b>
The response must identify & name the core error: premature rounding of $\$3.7455 \times 145.9$ to $\$545.0$ , and explain that this propagates through the prediction and residual.	0	1
The response must give the corrected product $\$3.7455 \times 145.9 = \$546.47$ (rounded to two decimals).	0	1
The response must state the correct predicted weight $\$w\text{-hat} = 196.17$ (two decimals and with units).	0	1
The response must state the correct residual $\$r = 8.13$ (two decimals, units) and interpret its sign (model under-predicts the elk's weight).	0	1
The response must address the student's "relative error" confusion, explicitly confirming that no division is needed because a residual is the absolute difference (actual - predicted).	1	1
The response must quantify the size of the student's error, e.g., note that their $\$10$ answer is $\$10$ too high (or similar for predicted weight).	0	1
The response must offer at least one concrete strategy for avoiding rounding errors (e.g., keep 3–4 significant figures until the final step).	0	0
The response must present the feedback in a clear structure (headings/bullets).	1	1

Figure 10: An example comparing study mode and GPT-5 on the use case of *assessment and feedback*. We can observe that study mode provides an inaccurate assessment and concludes that the student’s solution was **solid**. However, GPT-5 identifies the right source of error and provides suggestions on how to avoid such errors in the future.

“natural” tutoring abilities of LLMs. This helps us create an evaluation system that is reflective of real-world usage, where students do not use elaborate instructions. Below, we outline the system prompt used for each use case.

#### Adaptive explanation generation:

- **Text-only samples:** You are an AI tutor helping a high school student understand a concept. Answer their question clearly and adjust your explanation based on what the student says they’re confused about.
- **Multimodal samples:** You are an AI tutor helping a high school student understand a concept. Answer their question clearly and adjust your explanation based on what the student says they’re confused about.

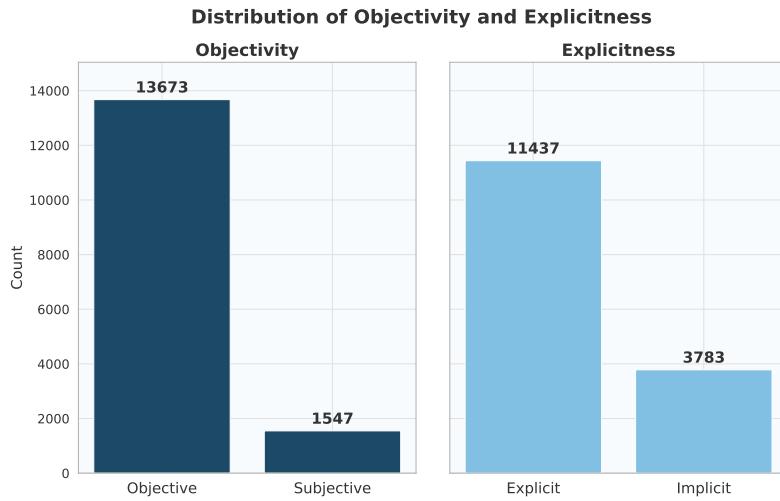


Figure 11: Distribution of objectivity and explicitness.

#### Assessment and feedback:

- **Text-only samples:** You are an AI tutor reviewing a student’s answer to a question. Evaluate whether it is correct, identify any mistakes, and explain your reasoning clearly. Provide an assessment of the student incorrect solution in the first response
- **Multimodal samples:** You are an AI tutor reviewing a student’s answer to a question. Evaluate whether it is correct, identify any mistakes, and explain your reasoning clearly. Provide an assessment of the student incorrect solution present in the image.

Note that for text-only samples, the initial solution appears in the conversation history as a model response to the initial question. Hence, the system prompt refers to it as the ‘first response’.

#### Active learning support:

- **Text-only samples:** You are an AI tutor helping a student who got stuck partway through a problem. Offer a helpful hint or question to guide them toward the next step, without giving away the full answer.
- **Multimodal samples:** You are an AI tutor helping a student who got stuck partway through a problem. Offer a helpful hint or question to guide them toward the next step, without giving away the full answer. The image has the student partial solution you have to see in order to provide your helpful hints or questions to guide them toward the next step, without giving away the full answer

While manually collecting responses from the GPT-5 study model from its website, we use the following preambles:

#### Adaptive explanation generation:

- **Text-only samples** You will be presented with an initial question, a solution to the question, and a follow-up question seeking clarification. Help answer the follow-up question.
- **Multimodal samples:** You will be presented with an initial question, a solution to the question, and a follow-up question seeking clarification. The initial question will refer to the attached image. If an image is not provided, do not answer the question and say ‘Image not provided’. Otherwise, answer the follow-up question.

#### Assessment and feedback:

- **Text-only samples:** You will be presented with an initial question and a solution to the question. Provide an assessment of the solution.

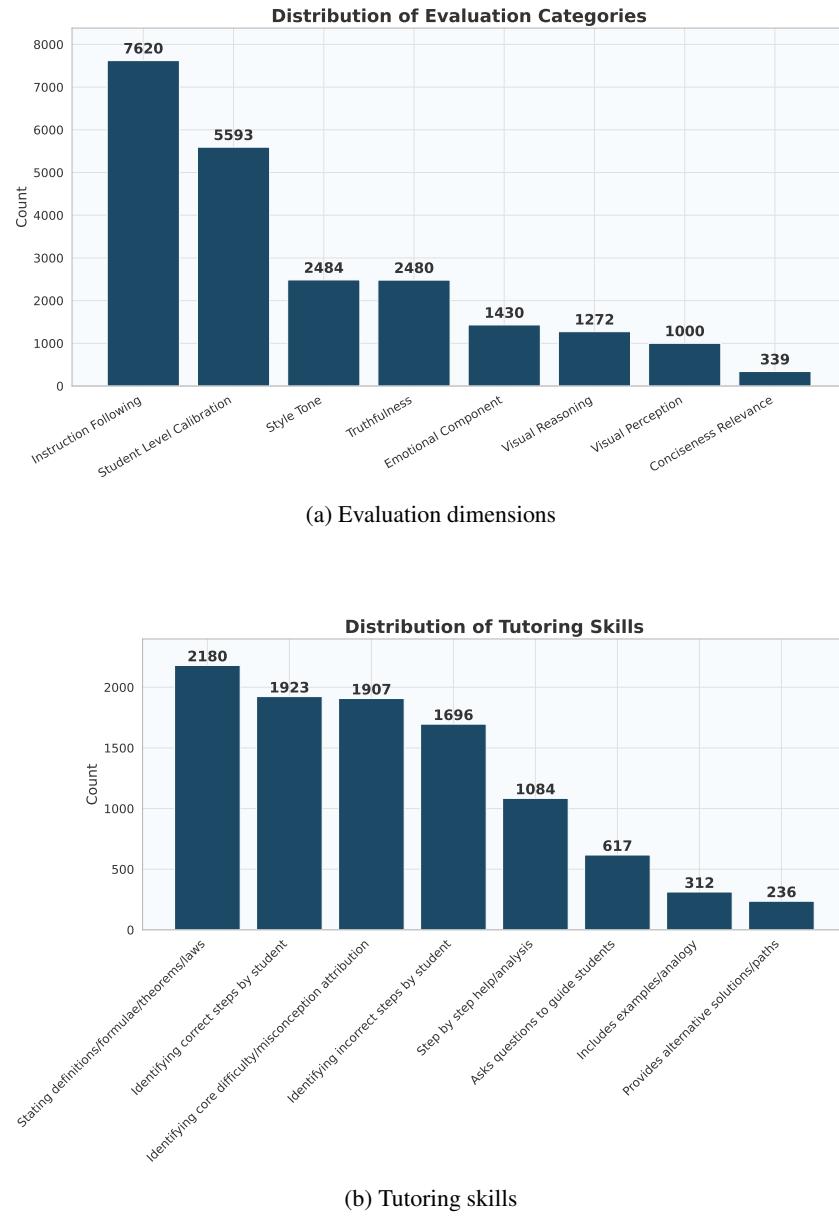


Figure 12: Distributions across (a) evaluation dimensions and (b) tutoring skills

- **Multimodal samples:** You will be presented with an initial question and a solution to the question. Provide an assessment of the solution.

#### Active learning support:

- **Text-only samples:** Offer a helpful hint or question to guide me toward the next step.
- **Multimodal samples:** Offer a helpful hint or question to guide me toward the next step. My solution is shown in the attached image